



**PG BootCamp Russia 2025 Ekaterinburg**

[PGBootCamp.ru](http://PGBootCamp.ru)

# **Работа с запросами с точки зрения DBA**

Александр Никитин. DBA Team.

# Немного обо мне


- С базами начал работать где-то в 2001 г. (MS FoxPro, FireBird, Oracle, PostgreSQL).
- DBA с 2014 года.
- С 2020 г. выступаю на крупных конференциях.
- В консалтинге с 2021 года.

# Содержание

- Инструменты
- Примеры использования
- Выводы

# Инструменты

## ЧТО И КАКИМ НОЖОМ РЕЗАТЬ



Людей, которые слушают музыку в общественных местах без наушников

# Инструменты

pg\_stat\_statements – расширение (что это?) PostgreSQL


```
psql: alter system set shared_preload_libraries='pg_stat_statements';
```

```
psql: create extension pg_stat_statements;
```

Рестарт PostgreSQL

<https://www.postgresql.org/docs/17/pgstatstatements.html>

<https://postgrespro.ru/education/books/monitoring>



# Инструменты

<https://github.com/dataegret/pg-utils> - набор утилит DBA

```
git clone https://github.com/dataegret/pg-utils ~/stuff
```



УТИЛИТЫ



Доклад

# Инструменты

```
psql -f ~/stuff/sql/globals/query_stat_total_N.sql
```

---

```
total time:      25:42:36 (IO: 0.99%)  
total queries:  77,966,682 (unique: 2,103)  
report for all databases, version 0.9.5 @ PostgreSQL 11.7  
tracking top 10000 queries, utilities on, logging 50ms+ queries
```

# Инструменты

```
psql -f ~/stuff/sql/globals/query_stat_total_N.sql
```

```
-----  
total time:      25:42:36 (IO: 0.99%)  
total queries:   77,966,682 (unique: 2,103)  
report for all databases, version 0.9.5 @ PostgreSQL 11.7  
tracking top 10000 queries, utilities on, logging 50ms+ queries
```

```
=====
```

pos:1	total time: 11:07:00 (43.2%, CPU: 43.7%, IO: 0.0%)	calls: 4,434 (0.01%)
-------	--	----------------------

```
  avg_time: 9025.74ms (IO: 0.0%)  
user: user1      db: db_name    rows: 12,817 (0.00%)    query:  
SELECT tote_location, COUNT(tote_location) FROM (SELECT * FROM (SELECT B.tote_name,  
B.tote_state, (CASE B.tote_location  
    WHEN $1...
```



# Инструменты

```
2024-09-29 00:00:00.686 MSK 105770 user@db_name from 10.13.19.81 [vxid:53/174763983
txid:0] [SELECT] LOG: duration: 100.494 ms execute S_17: select totetransp0_.id as
id1_1_, totetransp0_.created as created2_1_, totetransp0_.modified as modified3_1_,
totetransp0_.version as version4_1_, totetransp0_.active_date as active_d5_1_,
totetransp0_.complete_status as complete6_1_, totetransp0_.current_location_id as
current15_1_, totetransp0_.destination_location_id as destina16_1_,
totetransp0_.error_reason as error_re7_1_, totetransp0_.finished_date as finished8_1_,
totetransp0_.forcedfinish as forcedfi9_1_, totetransp0_.number as number10_1_,
totetransp0_.original_destination_location_id as original17_1_,
totetransp0_.source_location_id as source_18_1_, totetransp0_.status as status11_1_,
totetransp0_.tote_number as tote_nu12_1_, totetransp0_.transport_batch_number as
transpo13_1_, totetransp0_.type as type14_1_ from tote_transport_order totetransp0_ where
totetransp0_.tote_number=$1 and (totetransp0_.status in ($2 , $3 , $4 , $5 , $6)) order
by totetransp0_.modified desc limit $7
2024-09-29 00:00:00.686 MSK 105770 user@db_name from 10.13.19.81 [vxid:53/174763983
txid:0] [SELECT] DETAIL: parameters: $1 = 'T00000005212', $2 = 'CREATED', $3 = 'ACTIVE',
$4 = 'FINISHED', $5 = 'CANCELLED', $6 = 'ERROR', $7 = '1'
```

# Инструменты

```
2024-09-29 00:00:00.686 MSK 105770 user@db_name from 10.13.19.81 [vxid:53/174763983
txid:0] [SELECT] LOG: duration: 100.494 ms execute S_17: select totetransp0_.id as
id1_1_, totetransp0_.created as created2_1_, totetransp0_.modified as modified3_1_,
totetransp0_.version as version4_1_, totetransp0_.active_date as active_d5_1_,
totetransp0_.complete_status as complete6_1_, totetransp0_.current_location_id as
current15_1_, totetransp0_.destination_location_id as destina16_1_,
totetransp0_.error_reason as error_re7_1_, totetransp0_.finished_date as finished8_1_,
totetransp0_.forcedfinish as forcedfi9_1_, totetransp0_.number as number10_1_,
totetransp0_.original_destination_location_id as original17_1_,
totetransp0_.source_location_id as source_18_1_, totetransp0_.status as status11_1_,
totetransp0_.tote_number as tote_nu12_1_, totetransp0_.transport_batch_number as
transpo13_1_, totetransp0_.type as type14_1_ from tote_transport_order totetransp0_ where
totetransp0_.tote_number=$1 and (totetransp0_.status in ($2 , $3 , $4 , $5 , $6)) order
by totetransp0_.modified desc limit $7
2024-09-29 00:00:00.686 MSK 105770 user@db_name from 10.13.19.81 [vxid:53/174763983
txid:0] [SELECT] DETAIL: parameters: $1 = 'T00000005212', $2 = 'CREATED', $3 = 'ACTIVE',
$4 = 'FINISHED', $5 = 'CANCELLED', $6 = 'ERROR', $7 = '1'
```

# Инструменты

[https://github.com/Nikitin-Alexandr/utils/blob/main/replace\\_parameters](https://github.com/Nikitin-Alexandr/utils/blob/main/replace_parameters)



# Инструменты

```
CREATE OR REPLACE FUNCTION replace_parameters (query text, parameters
    text) RETURNS TEXT
    AS $body$
DECLARE
    rec
    record;
    SQL
    text;
BEGIN
    SQL := '';
    FOR rec IN
        ( SELECT
            array_to_string(REGEXP_MATCHES(parameters, '(<=\$)(\d+)(?=\ =\ )', 'g'), ';')
            AS arg_no, array_to_string(REGEXP_MATCHES(parameters, '(<=\ =\ )(.*?)(?=(,\ \
            \$\d+|\$))', 'g'), ';') AS arg_val)
        LOOP
            SQL:=regexp_replace(query, '\$'||rec.arg_no||'(!\d)',
                rec.arg_val,'g'); query := sql;
        END LOOP;
    return
    sql; END;
$body$
LANGUAGE 'plpgsql';
```

```
/*Example: select replace_parameters($$query$$,$$parameters$$);*/
```

# Примеры

- pg\_stat\_statements - расширение  
(<https://www.postgresql.org/docs/17/pgstatstatements.html>)
- <https://github.com/dataegret/pg-utils> - набор утилит DBA
- [https://github.com/Nikitin-Alexandr/Utils/blob/main/replace\\_parameters](https://github.com/Nikitin-Alexandr/Utils/blob/main/replace_parameters)  
замена параметров

# Примеры

## Создание индекса



# Примеры

```
select "id", "external_id", "party_id", "shop_id", "invoice_id", "payment_id"
from "payment"
where "status_created_at" >= cast('2018-09-30 21:00:00' as timestamp) and
      "status_created_at" < cast('2018-10-31 21:00:00' as timestamp) and
      "party_id" = '58cad298' and
      "shop_id" = 'de788f93' and
      "status" = 'captured'
order by "status_created_at" desc;
```

# Примеры


```
select "id", "external_id", "party_id", "shop_id", "invoice_id", "payment_id"
from "payment"
where "status_created_at" >= cast('2018-09-30 21:00:00' as timestamp) and
      "status_created_at" < cast('2018-10-31 21:00:00' as timestamp) and
      "party_id" = '58cad298' and
      "shop_id" = 'de788f93' and
      "status" = 'captured'
order by "status_created_at" desc;
```



# Примеры

```
select "id", "external_id", "party_id", "shop_id", "invoice_id", "payment_id"  
from "payment"  
where "status_created_at" >= cast('2018-09-30 21:00:00' as timestamp) and  
      "status_created_at" < cast('2018-10-31 21:00:00' as timestamp) and  
      "party_id" = '58cad298' and  
      "shop_id" = 'de788f93' and  
      "status" = 'captured'  
order by "status_created_at" desc;
```

Какие ваши доказательства предложения?



# Примеры

```
SELECT
  count(*) FILTER (WHERE party_id = '58cad298') cnt_party,
  count(*) FILTER (WHERE shop_id = 'de788f93') cnt_shop,
  count(*) FILTER (WHERE party_id = '58cad298' AND shop_id = 'de788f93') cnt_party_shop,
  count(*) FILTER (WHERE status_created_at >= '2020-12-31 21:00:00' AND status_created_at < '2021-04-09 21:00:00')
cnt_st_created,
  count(*) FILTER (WHERE status_created_at >= '2020-12-31 21:00:00' AND status_created_at < '2021-04-09 21:00:00' and
    "status" = 'captured') cnt_st_created_status,
  count(*) FILTER (WHERE shop_id = 'de788f93' AND status_created_at >= '2020-12-31 21:00:00' AND
    status_created_at < '2021-04-09 21:00:00') cnt_shop_created,
  count(*) FILTER (WHERE party_id = '58cad298' AND status_created_at >= '2020-12-31 21:00:00' AND
    status_created_at < '2021-04-09 21:00:00') cnt_party_created,
  count(*) FILTER (WHERE party_id = '58cad298' AND shop_id = 'de788f93' AND status_created_at >= '2020-12-31 21:00:00'
    AND status_created_at < '2021-04-09 21:00:00') all_filter
FROM payment \gx
```

# Примеры

cnt_party		329 810	(rows)
cnt_shop		326 743	
cnt_party_shop		326 743	
cnt_st_created		1 112 292	
cnt_st_created_status		910 527	
cnt_shop_created		10 635	
cnt_party_created		10 783	
all_filter		9 212	

# Примеры

## Статистика

Представление: `pg_stats`

Параметр: `default_statistics_target`

Поля: `null_frac`, `n_distinct`, `correlation`

300\*dst блоков и  
из них 300\*dst строк

`null_frac` – доля null значений

Если `n_distinct > 0` – это количество уникальных значений

Если `n_distinct < 0` – это доля уникальных значений

`Correlation` – корреляция между расположением данных на диске и однородностью их изменения. Если последующее всегда больше, предыдущего, то будет близко к 1. Если всегда меньше ближе к -1.

# Примеры

```
select null_frac, n_distinct, correlation from pg_stats where tablename =  
'payment' and attname = 'party_id' \gx
```

```
null_frac      | 0.000246667
```

```
n_distinct     | 498
```

```
correlation    | 0.0966761
```

```
select null_frac, n_distinct, correlation from pg_stats where tablename =  
'payment' and attname = 'shop_id' \gx
```

```
null_frac      | 0.000246667
```

```
n_distinct     | 1289
```

```
correlation    | 0.236482
```

# Примеры

- `count(*) FILTER (WHERE column_name = column_value)` — пример подсчёта
- `pg_stats (null_frac, n_distinct > 0, n_distinct < 0, correlation)` - представление со статистикой
- `default_statistics_target` - параметр (детализация статистики)

# Примеры

OR в условии запроса



# Примеры

```
select * from tickets_ticket, auth_user
where (auth_user.id = tickets_ticket.buyer_id) and
((upper((username)::text) = '406240'::text) OR (tickets_ticket.buyer_id =
406240));
```




# Примеры

```
select * from tickets_ticket, auth_user
where (auth_user.id = tickets_ticket.buyer_id) and
((upper((username)::text) = '406240'::text) OR (tickets_ticket.buyer_id =
406240));
```

# Примеры

```
select * from tickets_ticket, auth_user  
where (auth_user.id = tickets_ticket.buyer_id) and  
((upper((username)::text) = '406240'::text) OR (tickets_ticket.buyer_id =  
406240));
```

Execution Time: 814.906 ms




# Примеры

```
select * from tickets_ticket, auth_user
where (auth_user.id = tickets_ticket.buyer_id) and
((upper((username)::text) = '406240'::text) OR (tickets_ticket.buyer_id =
406240));
```

Execution Time: 814.906 ms

```
select * from tickets_ticket, auth_user
where auth_user.id = tickets_ticket.buyer_id and upper((username)::text) =
'406240'::text
union
select * from tickets_ticket, auth_user
where auth_user.id = tickets_ticket.buyer_id and tickets_ticket.buyer_id =
406240;
```




# Примеры

```
select * from tickets_ticket, auth_user
where (auth_user.id = tickets_ticket.buyer_id) and
((upper((username)::text) = '406240'::text) OR (tickets_ticket.buyer_id =
406240));
```

Execution Time: 814.906 ms

```
select * from tickets_ticket, auth_user
where auth_user.id = tickets_ticket.buyer_id and upper((username)::text) =
'406240'::text
union
select * from tickets_ticket, auth_user
where auth_user.id = tickets_ticket.buyer_id and tickets_ticket.buyer_id =
406240;
```

Execution Time: 76.741 ms



# Примеры

`(auth_user.id =  
tickets_ticket.buyer_id)`

`(upper((username)::text) = '406240'::text)`

`(tickets_ticket.buyer_id = 406240)`

# Примеры

```
select * from tickets_ticket,auth_user
where auth_user.id = tickets_ticket.buyer_id and upper((username)::text) =
'406240'::text and tickets_ticket.buyer_id != 406240
UNION ALL
select * from tickets_ticket,auth_user
where auth_user.id = tickets_ticket.buyer_id and tickets_ticket.buyer_id =
406240 and upper((username)::text) != '406240'::text
UNION ALL
select * from tickets_ticket,auth_user
where auth_user.id = tickets_ticket.buyer_id and tickets_ticket.buyer_id =
406240 and upper((username)::text) = '406240'::text;
Execution Time: 28.147 ms
```

# Примеры

```
select * from tickets_ticket,auth_user
where auth_user.id = tickets_ticket.buyer_id and upper((username)::text) =
'406240'::text and tickets_ticket.buyer_id != 406240
UNION ALL
select * from tickets_ticket,auth_user
where auth_user.id = tickets_ticket.buyer_id and tickets_ticket.buyer_id =
406240 and upper((username)::text) != '406240'::text
UNION ALL
select * from tickets_ticket,auth_user
where auth_user.id = tickets_ticket.buyer_id and tickets_ticket.buyer_id =
406240 and upper((username)::text) = '406240'::text;
Execution Time: 28.147 ms
```

До модификации:

Execution Time: 814.906 ms – ускорение ~ 30 раз.

# Примеры

- Запросы, содержащие OR, можно переписать на ряд запросов с использованием UNION ALL



# Примеры

## Distinct + join



# Примеры

```
select distinct tbl1.id  
from tbl1  
    inner join tbl2 on tbl1.id_pp = tbl2.id_pp;
```

# Примеры

```
select distinct tbl1.id  
from tbl1  
    inner join tbl2 on tbl1.id_pp = tbl2.id_pp;
```

```
select id  
from tbl1  
where exists (  
    select 1  
    from tbl2  
    where tbl1.id_pp = tbl2.id_pp);
```

<https://habr.com/ru/company/tensor/blog/513324/>



# Примеры

```
create table t1 (id int, t text);
insert into t1 values(1,'a'), (1,'a'), (2,'b'), (3,'c');
create table t2 (id int, val text);
insert into t2 values (1,'true'),(2,'true'),(3,'false');
select distinct t1.* from t1 inner join t2 on t1.id = t2.id and t2.val = 'true';
  1 | a
  2 | b
select t1.* from t1 where exists (select 1 from t2 where t1.id = t2.id and t2.val =
  'true');
  1 | a
  1 | a
  2 | b
```

# Примеры

Проверка эквивалентности запросов:

```
select * from (original_query) a  
except  
select * from (rewritten_query) b;
```

# Примеры

Проверка эквивалентности запросов:

```
select * from (original_query) a  
except  
select * from (rewritten_query) b;
```

```
select * from (rewritten_query) b  
except  
select * from (original_query) a;
```

# Примеры

- `distinct+join` — за это сочетание должен цепляться взгляд при анализе.

# Примеры

## Работа со статистикой





# Примеры

```
Index Scan using personal_grades_reference_id_key on personal_grades g (cost=0.42..2862542.35 rows=1 width=32) (actual time=572.884..572.898 rows=1 loops=1)
Index Cond: ((reader_id = 102611615013>4216-8005-91952c383f20::uuid) AND (reference_id = 59))
Buffers: shared hit=183 read=24 drect=1
IO Timings: read=0.452
SubPlan 4
-> Aggregate (cost=2862541.67..2862541.68 rows=1 width=32) (actual time=4.523..4.533 rows=1 loops=1)
  Buffers: shared hit=176 read=24
  IO Timings: read=0.452
  -> Subquery Scan on t1 (cost=2862541.50..2862541.62 rows=10 width=848) (actual time=3.730..3.742 rows=1 loops=1)
    Buffers: shared hit=176 read=24
    IO Timings: read=0.452
    -> Sort (cost=2862541.50..2862541.52 rows=10 width=856) (actual time=3.571..3.580 rows=1 loops=1)
      Sort Key: t1."position"
      Sort Method: quicksort Memory: 201kB
      Buffers: shared hit=170 read=24
      IO Timings: read=0.452
      -> Index Scan using personal_topics_grade_id_key on personal_topics t1 (cost=0.43..2862541.33 rows=10 width=856) (actual time=3.386..3.398 rows=1 loops=1)
        Index Cond: (grade_id = g.id)
        Filter: (deleted IS FALSE)
        Buffers: shared hit=170 read=24
        IO Timings: read=0.452
        SubPlan 3
        -> Aggregate (cost=286254.02..286254.03 rows=1 width=32) (actual time=2.937..2.945 rows=1 loops=1)
          Buffers: shared hit=170 read=23
          IO Timings: read=0.219
          -> Subquery Scan on ch (cost=286253.60..286253.90 rows=24 width=335) (actual time=2.610..2.638 rows=5 loops=1)
            Buffers: shared hit=170 read=23
            IO Timings: read=0.219
            -> Sort (cost=286253.60..286253.66 rows=24 width=343) (actual time=2.590..2.598 rows=5 loops=1)
              Sort Key: ch."position"
              Sort Method: quicksort Memory: 199kB
              Buffers: shared hit=170 read=23
              IO Timings: read=0.219
              -> Index Scan using personal_chapters_topic_id_key on personal_chapters ch_1 (cost=0.43..286253.05 rows=24 width=343) (actual time=0.658..2.525 rows=5 loops=1)
                Index Cond: (topic_id = t1.id)
                Filter: (deleted IS FALSE)
                Buffers: shared hit=170 read=23
                IO Timings: read=0.219
                SubPlan 2
                -> Aggregate (cost=12010.48..12010.49 rows=1 width=32) (actual time=0.483..0.484 rows=1 loops=5)
                  Buffers: shared hit=166 read=22
                  IO Timings: read=0.208
                  -> Subquery Scan on l1 (cost=12005.74..12009.12 rows=271 width=1579) (actual time=0.396..0.396 rows=5 loops=5)
                    Buffers: shared hit=166 read=22
                    IO Timings: read=0.208
                    -> Sort (cost=1009.74..12006.41 rows=271 width=1587) (actual time=0.380..0.381 rows=5 loops=5)
                      Sort Key: l1."position"
                      Sort Method: quicksort Memory: 57kB
                      Buffers: shared hit=166 read=22
                      IO Timings: read=0.208
                      -> Index Scan using personal_lessons_chapter_id_key on personal_lessons L1 (cost=0.56..11994.79 rows=271 width=1587) (actual time=0.110..0.366 rows=5 loops=5)
                        Index Cond: (chapter_id = ch_1.id)
                        Filter: (deleted IS FALSE)
                        Buffers: shared hit=166 read=22
                        IO Timings: read=0.208
                        SubPlan 1
                        -> Aggregate (cost=44.20..44.21 rows=1 width=32) (actual time=0.061..0.061 rows=1 loops=27)
                          Buffers: shared hit=134 read=21
                          IO Timings: read=0.192
                          -> Subquery Scan on c (cost=36.90..42.11 rows=417 width=715) (actual time=0.018..0.022 rows=4 loops=27)
                            Buffers: shared hit=134 read=21
                            IO Timings: read=0.192
                            -> Sort (cost=36.90..37.94 rows=417 width=691) (actual time=0.016..0.017 rows=4 loops=27)
                              Sort Key: c."position"
                              Sort Method: quicksort Memory: 29kB
                              Buffers: shared hit=134 read=21
                              IO Timings: read=0.192
                              -> Index Scan using personal_cards_lesson_id_key on personal_cards c_1 (cost=0.56..18.75 rows=417 width=691) (actual time=0.008..0.014 rows=4 loops=27)
                                Index Cond: (lesson_id = L1.id)
                                Filter: (deleted IS FALSE)
                                Buffers: shared hit=134 read=21
                                IO Timings: read=0.192
Planning Time: 3.397 ms
JIT:
Functions: 90
Options: infining true, Optimization true, Expressions true, Deforming true
Timing: Generation 0.004 ms, Infining 105.226 ms, Optimization 290.538 ms, Emission 170.609 ms, Total 573.036 ms
Execution Time: 616.628 ms
(78 rows)
```

# Примеры

```
Index Scan using personal_grades_reference_id_key on personal_grades g (cost=0.42..2862542.35 rows=1 width=32) (actual time=572.884..572.898 rows=1 loops=1)
Index Cond: ((reader_id = 1026116f5d3d3422ba4b0c591952c383f20::uuid) AND (reference_id = 59))
Buffers: shared hit=183 read=24 drect=1
IO Timings: read=0.452
SubPlan 4
-> Aggregate (cost=2862541.67..2862541.68 rows=1 width=32) (actual time=4.523..4.533 rows=1 loops=1)
  Buffers: shared hit=176 read=24
  IO Timings: read=0.452
  -> Subquery Scan on t1 (cost=2862541.50..2862541.62 rows=10 width=848) (actual time=3.730..3.742 rows=1 loops=1)
    Buffers: shared hit=176 read=24
    IO Timings: read=0.452
    -> Sort (cost=2862541.50..2862541.52 rows=10 width=856) (actual time=3.571..3.580 rows=1 loops=1)
      Sort Key: t1."position"
      Sort Method: quicksort Memory: 201kB
      Buffers: shared hit=170 read=24
      IO Timings: read=0.452
      -> Index Scan using personal_topics_grade_id_key on personal_topics t1_1 (cost=0.43..2862541.33 rows=10 width=856) (actual time=3.386..3.398 rows=1 loops=1)
        Index Cond: (grade_id = g.id)
        Index Cond: (deleted IS FALSE)
        Buffers: shared hit=170 read=24
        IO Timings: read=0.452
        SubPlan 3
        -> Aggregate (cost=286254.02..286254.03 rows=1 width=32) (actual time=2.937..2.945 rows=1 loops=1)
          Buffers: shared hit=170 read=23
          IO Timings: read=0.219
          -> Subquery Scan on ch (cost=286253.60..286253.90 rows=24 width=335) (actual time=2.610..2.638 rows=5 loops=1)
            Buffers: shared hit=170 read=23
            IO Timings: read=0.219
            -> Sort (cost=286253.60..286253.66 rows=24 width=343) (actual time=2.590..2.598 rows=5 loops=1)
              Sort Key: ch."position"
              Sort Method: quicksort Memory: 199kB
              Buffers: shared hit=170 read=23
              IO Timings: read=0.219
              -> Index Scan using personal_chapters_topic_id_key on personal_chapters ch_1 (cost=0.43..286253.05 rows=24 width=343) (actual time=0.658..2.525 rows=5 loops=1)
                Index Cond: (topic_id = t1.id)
                Index Cond: (deleted IS FALSE)
                Buffers: shared hit=170 read=23
                IO Timings: read=0.219
                SubPlan 2
                -> Aggregate (cost=12010.48..12010.49 rows=1 width=32) (actual time=0.483..0.484 rows=1 loops=5)
                  Buffers: shared hit=166 read=22
                  IO Timings: read=0.208
                  -> Subquery Scan on l1 (cost=12005.74..12009.12 rows=271 width=1579) (actual time=0.386..0.395 rows=5 loops=5)
                    Buffers: shared hit=166 read=22
                    IO Timings: read=0.208
                    -> Sort (cost=1006.75..12006.41 rows=271 width=1587) (actual time=0.380..0.381 rows=5 loops=5)
                      Sort Key: l1."position"
                      Sort Method: quicksort Memory: 57kB
                      Buffers: shared hit=166 read=22
                      IO Timings: read=0.208
                      -> Index Scan using personal_lessons_chapter_id_key on personal_lessons l1_1 (cost=0.56..11994.79 rows=271 width=1587) (actual time=0.110..0.366 rows=5 loops=5)
                        Index Cond: (chapter_id = ch_1.id)
                        Index Cond: (deleted IS FALSE)
                        Buffers: shared hit=166 read=22
                        IO Timings: read=0.208
                        SubPlan 1
                        -> Aggregate (cost=44.20..44.21 rows=1 width=32) (actual time=0.061..0.061 rows=1 loops=27)
                          Buffers: shared hit=134 read=21
                          IO Timings: read=0.192
                          -> Subquery Scan on c (cost=36.90..42.11 rows=417 width=715) (actual time=0.018..0.022 rows=4 loops=27)
                            Buffers: shared hit=134 read=21
                            IO Timings: read=0.192
                            -> Sort (cost=36.90..37.94 rows=417 width=691) (actual time=0.016..0.017 rows=4 loops=27)
                              Sort Key: c."position"
                              Sort Method: quicksort Memory: 29kB
                              Buffers: shared hit=134 read=21
                              IO Timings: read=0.192
                              -> Index Scan using personal_cards_lesson_id_key on personal_cards c_1 (cost=0.56..18.75 rows=417 width=691) (actual time=0.008..0.014 rows=4 loops=27)
                                Index Cond: (lesson_id = l1.id)
                                Index Cond: (deleted IS FALSE)
                                Buffers: shared hit=134 read=21
                                IO Timings: read=0.192
Planning Time: 3.397 ms
JIT:
Functions: 90
Options: infining true, Optimization true, Expressions true, Deforming true
Timing: Generation 0.004 ms, Infining 105.226 ms, Optimization 290.538 ms, Emission 170.609 ms, Meta 573.036 ms
Execution Time: 616.628 ms
(78 rows)
```

# Примеры

-> Index Scan using personal\_cards\_lesson\_id\_key on personal\_cards c\_1  
(cost=0.56..18.75 rows=417 width=691) (actual time=0.008..0.014 rows=4  
loops=27)

Index Cond: (lesson\_id = l\_1.id) Filter:  
(deleted IS FALSE) Buffers: shared  
hit=134 read=21 I/O Timings: read=0.192

Planning Time: 3.397 ms JIT:

Functions: 50

Options: Inlining true, Optimization true, Expressions true, Deforming true

Timing: Generation 6.664 ms, Inlining 105.226 ms, Optimization 290.538 ms, Emission  
170.609 ms,

Total 573.036 ms Execution Time:

616.528 ms

# Примеры

```
->      Index Scan using personal_cards_lesson_id_key on personal_cards c_1  
      (cost=0.56..18.75 rows=417 width=691) (actual time=0.008..0.014 rows=4 loops=27)
```

```
Index Cond: (lesson_id = l_1.id)
```

```
Filter: (deleted IS FALSE) Buffers:
```

```
shared hit=134 read=21 I/O Timings:
```

```
read=0.192
```

```
Planning Time: 3.397 ms
```

```
JIT:
```

```
Functions: 50
```

```
Options: Inlining true, Optimization true, Expressions true, Deforming true
```

```
Timing: Generation 6.664 ms, Inlining 105.226 ms, Optimization 290.538 ms, Emission 170.609 ms,
```

```
Total 573.036 ms Execution
```

```
Time: 616.528 ms
```

В PostgreSQL тоже есть JIT :)

# Примеры

```
analyze personal_cards;  
explain select distinct lesson_id from personal_cards;  
QUERY PLAN
```

---

```
Unique  (cost=0.56..1257422.12 rows=151734 width=8)  
->  Index Only Scan using personal_cards_lesson_id_key on personal_cards  
(cost=0.56..1092318.52 rows=66041439 width=8)
```

```
select count(*), count(distinct lesson_id) from personal_cards;  
count      | count  
-----+-----  
65844906 | 22326223
```

# Примеры

```
analyze personal_cards;  
explain select distinct lesson_id from personal_cards;  
QUERY PLAN
```

---

```
Unique (cost=0.56..1257422.12 rows=151734 width=8)  
-> Index Only Scan using personal_cards_lesson_id_key on personal_cards  
(cost=0.56..1092318.52 rows=66041439 width=8)
```

```
select count(*), count(distinct lesson_id) from personal_cards;  
count      | count  
-----+-----  
65844906 | 22326223
```

```
select count(*), count(distinct lesson_id) from personal_cards tablesample system (N);  
N - проценты от всей таблицы
```

# Примеры

```
select count(*), count(distinct lesson_id) from personal_cards;
```

count		count
-----+-----		
65844906		22326223

```
select 22326223::float/65844906;  
?column?
```

```
-----  
0.3390728965426726
```

```
alter table personal_cards alter lesson_id set (n_distinct=-0.4);
```

```
analyze personal_cards;
```

# Примеры

```
select count(*), count(distinct lesson_id) from personal_cards;
```

count		count
-----+-----		
65844906		22326223

```
select 22326223::float/65844906;  
?column?
```

```
-----  
0.3390728965426726
```

```
alter table personal_cards alter lesson_id set (n_distinct=-0.4);
```

```
analyze personal_cards;
```

**616 MC -> 4 MC**



# Примеры

- Мы можем переписать значение `n_distinct` исходя из наших представлений.
- Помним о том, что представляют из себя положительные и отрицательные значения в `n_distinct`.

# Примеры

Пытаемся понять, что имел в виду автор запроса




# Примеры

```
SELECT "id", "jdbc_id", ... FROM (  
  SELECT  
    t.id as jdbc_id,  
    ...  
  FROM  
    tbl_1 as t  
  WHERE created_at BETWEEN $5::date - interval $6 day AND $7::date) as x  
WHERE "jdbc_id" >= $8 AND "jdbc_id" < $9;
```

# Примеры

```
SELECT "id", "jdbc_id", ... FROM (  
  SELECT  
    t.id as jdbc_id,  
    ...  
  FROM  
    tbl_1 as t  
  WHERE created_at BETWEEN $5::date - interval $6 day AND $7::date) as x  
WHERE "jdbc_id" >= $8 AND "jdbc_id" < $9;
```



# Примеры

- Поймите, чего именно хотел добиться автор запроса, когда писал его.
- Подумайте, может быть вы можете решить эту задачу лучше.

# Примеры

## Сложный запрос



# Примеры

```
SELECT COUNT(*)
FROM "_1_"
  INNER JOIN "_2_" ON "_2_". "id" = "_1_". "clb_id"
  INNER JOIN "_3_" ON "_3_". "id" = "_2_". "t_id"
WHERE "_3_". "s_id" IN
  (SELECT "_3_". "s_id"
   FROM "_1_"
     INNER JOIN "_2_" ON "_2_". "id" = "_1_". "clb_id"
     INNER JOIN "_3_" ON "_3_". "id" = "_2_". "t_id"
     INNER JOIN "_4_" ON "_1_". "id" = "_4_". "tmp_id"
   WHERE "_4_". "st_id" = 6251961 AND "_4_". "deleted_at" IS NULL)
  AND "_1_". "m_id" = 192 AND "_1_". "clb_id" IN
  (SELECT "_2_". "id"
   FROM "_2_"
   WHERE "_2_". "t_id" IN
    (SELECT "_3_". "id" FROM "_3_" WHERE "_3_". "s_id" = 262718)) AND ("_1_". "points" > 0)
```

# Примеры

**from\_collapse\_limit** - Задаёт максимальное число элементов в списке FROM, до которого планировщик будет объединять вложенные запросы с внешним запросом.

При меньших значениях сокращается время планирования, но план запроса может стать менее эффективным.

**join\_collapse\_limit** - Задаёт максимальное количество элементов в списке FROM, до достижения которого планировщик будет сносить в него явные конструкции JOIN (за исключением FULL JOIN).

При меньших значениях сокращается время планирования, но план запроса может стать менее эффективным.



# Примеры

```
from_collapse_limit/join_collapse_limit    = 8  
Execution Time: 640.367 ms
```

# Примеры

```
from_collapse_limit/join_collapse_limit    = 8  
Execution Time: 640.367 ms
```

```
from_collapse_limit/join_collapse_limit    = 10  
Execution Time: 2.675 ms
```

<https://habr.com/ru/company/postgrespro/blog/574702/>



# Примеры



- `from_collapse_limit` и `join_collapse_limit` - опасные параметры.

# Примеры

## Американские горки



# Примеры



# Примеры

- Бывают ситуации, когда ДБА может вам только указать на то, что есть проблема, но найти и устранить причину иногда можете только вы.

# Выводы

1) Ваши друзья не только индексы.

# Выводы

- 1) Ваши друзья не только индексы.
- 2) Не стесняйтесь смотреть в статистику.



# Выводы



Чем ты занимаешься после работы?



Я PostgreSQL DBA...



- 1) Ваши друзья не только индексы.
- 2) Не стесняйтесь смотреть в статистику.
- 3) Любите ваших ДБА :)

# DBA Team

**Ваши данные - наша забота!**

- **Поддержка PostgreSQL 24/7**
- **Настройка кластеров конкретно под ваш профиль нагрузки;**
- **Помощь в оптимизации SQL-запросов;**
- **Выполнение миграций данных без простоя системы;**
- **Миграция БД на новые сервера или в другой ДЦ с минимальным простоем;**
- **Своевременное обновление PostgreSQL (как минорное, так и мажорное);**
- **Настройка резервного копирования и восстановления;**
- **Консультации по архитектуре и производительности БД;**
- **Проведение регулярных проверок ваших кластеров;**



**PG BootCamp Russia 2025 Ekaterinburg**

[PGBootCamp.ru](https://pgbootcamp.ru)

**Вопросы?**

Александр Никитин

E-mail: [contact@dba.team](mailto:contact@dba.team)

tg: @anikitindba

Бонус: знаем команду DevOps.